

# Operating Systems Lab

## Dept of CSE, LNMIIT, Jaipur

### 01 Feb 2021

### Lab Session 03

This document is based on an IIIT, Guwahati lab exercise. This assignment has four goals:

1. To dust off your C/C++ programming skills;
2. To observe the operating systems' (OS') activities at the user-level;
3. To implement an OS command interpreter (shell);
4. To use some system calls (such as `fork()` and `exec()`).

#### Part I: Observing the OS through the `/proc` file system

The OS is a program that uses various data structures. Like all programs in execution, you can determine the performance and other behaviour of the OS by inspecting its state - the values stored in its data structures. In this part of the assignment, we study some aspects of the organization and behaviour of a Linux system by observing values of kernel data structures exposed through the `/proc` virtual file system.

##### The `/proc` virtual file system:

Linux uses `/proc` file system to collect information from kernel data structures. The `/proc` implementation provided with Linux can read many different kernel data structures. If you `cd` to `/proc` on a Linux machine, you will see a number of files and directories at that location. Each file in this directory subtree corresponds to some kernel data structure. The subdirectories with numeric names contain virtual files with information about the process whose process ID matches the directory name. Files in `/proc` can be read like ordinary ASCII files. You can open each file and read it using library routines such as `fgets()` or `fscanf()`. The `proc(5)` manual page explains the virtual files and their content available through the `/proc` file system.

##### Requirements in detail:

In this part, you are asked to write a program to report the behaviour of the Linux kernel. Your program should run in two different versions. The default version should print the following values on `stdout`:

- Processor type
- Kernel version
- The amount of memory configured into this computer
- Amount of time since the system was last booted

A second version of the program should [run continuously](#) and print lists of the following dynamic values (each value in the lists is the average over a specified interval):

- The percentage of time the processor(s) spend in user mode, system mode, and the percentage of time the processor(s) are idle

- The amount and percentage of available (or free) memory
- The rate (number of sectors per second) of disk read/write in the system
- The rate (number per second) of context switches in the kernel
- The rate (number per second) of process creations in the system

If your program (compiled executable) is called `proc_parse`, running it without any parameter should print out information required for the first version. Running it with two parameters "`proc_parse <read_rate> <printout_rate>`" should print out information required for the second version. `read_rate` represents the time interval between two consecutive reads on the `/proc` file system. `printout_rate` indicates the time interval over which the average values should be calculated. Both `read_rate` and `printout_rate` are in seconds. For instance, `proc_parse 2 60` should read kernel data structures once every two seconds. It should then print out averaged kernel statistics once a minute (average of 30 samples). The second version of your program doesn't need to terminate.

## Part II: Building a shell

### UNIX shells:

The OS command interpreter is a program that people interact with to launch and control programs. On UNIX systems, the command interpreter is often called **shell**: a user-level program that gives people a command-line interface to launching, suspending, and killing other programs. `sh`, `ksh`, `csh`, `tcsh`, `bash`, ... are all examples of UNIX shells. You use a shell like this every time you log into a Linux machine and bring up a terminal. It might be useful to look at the manual pages of these shells, for example, type "`man csh`". Bourne shell `bsh` is commonly used shell at LNMIIT labs.

The most rudimentary shell is structured as the following loop:

1. Print out a prompt (e.g., "`CSC2/456Shell$` ");
2. Read a line from the user;
3. Parse the line into the program name and an array of parameters.
4. Use the `fork()` system call to spawn a new child process;
  - The child process then uses the `exec()` system call (or one of its variants) to launch the specified program;
  - The parent process (the shell) uses the `wait()` system call (or one of its variants) to wait for the child to terminate;
5. Once the child (the launched program) finishes, the shell repeats the loop by jumping to 1.

Although most commands people type on the shell prompt are the names of other UNIX programs (such as `ps` or `cat`), shells also recognize some special commands (called internal commands) that are not program names. For example, the `exit` command terminates the shell, and the `cd` command changes the current working directory. Shells directly make system calls to execute these commands, instead of forking a child process to handle them.

### Requirements in detail:

Your job is to implement a very primitive shell that knows how to launch new programs in the foreground and the background. It should also recognize a few internal commands. More specifically, it should support the following features.

- It should recognize the internal commands: `exit`, `source`, and `cd`. `exit` should use the `exit()` system call to terminate the shell. `cd` uses the `chdir()` system call to change to a new directory.
- If the command line does not indicate any internal commands, it should be in the following form:  
`<program name> <arg1> <arg2> .... <argN> [&]`  
 Your shell should invoke the program, passing it the list of arguments in the command line. The shell must wait until the started program completes unless the user runs it in the background (with `&`).

To allow users to pass arguments you need to parse the input line into words separated by whitespace (spaces and `\t` tab characters). You might try to use `strtok_r()` for parsing (check the manual page of `strtok_r()` and Google it for examples of using it). In case you wonder, `strtok_r()` is a user-level utility, not a system call. This means this function is fulfilled without the help of the operating system kernel. To make the parsing easy for you, you can assume the `'&'` token (when used) is separated from the last argument with one or more spaces or `\t` tab characters.

The shell runs programs using two core system calls: `fork()` and `execvp()`. Read the manual pages to see how to use them. In short, `fork()` creates an exact copy of the currently running process, and is used by the shell to spawn a new process. The `execvp()` call is used to overload the currently running program with a new program, which is how the shell turns a forked process into the program it wants to run. In addition, the shell must wait until the previously started program completes unless the user runs it in the background (with `&`). This is done with the `wait()` system call or one of its variants (such as `waitpid()`). All these system calls can fail due to unforeseen reasons (see their manual pages for details). You should check their return status and report errors if they occur.

No input the user gives should cause the shell to exit (except when the user types `exit` or `Ctrl+D`). This means your shell should handle errors gracefully, no matter where they occur. Even if an error occurs in the middle of a long pipeline, it should be reported accurately and your shell should recover gracefully. In addition, your shell should not generate leaking open file descriptors. **Hint:** you can monitor the current open file descriptors of the shell process through the `/proc` file system. Use C library functions `setjmp()` and `longjmp()` in `setjmp.h` if needed. (<https://en.wikipedia.org/wiki/Setjmp.h>)

### Pipes:

Your shell needs to support pipes. Pipes allow the `stdins` and `stdouts` of a list of programs to be concatenated in a chain. More specifically, the first program's `stdout` is directed to the `stdin` of the second program; the second program's `stdout` is directed to the `stdin` of the third program; and so on so forth. Multiple piped programs in a command line are separated with the token `"|"`. A command line will therefore have the following form:

`<program1> <arglist1> | <program2> <arglist2> | ... | <programN> <arglistN> [&]`

Try an example like this: pick a text file with more than 10 lines (assume it is called `textfile`) and then type

```
cat textfile | gzip -c | gunzip -c | tail -n 10
```

in a regular shell. Pause a bit to think what it really does. Note that multiple processes need to be launched for piped commands and all of them should be waited on in a foreground execution. The `pipe()` and `dup2()` system calls will be useful.